

Contents

1	List decoding of Reed-Solomon codes	1
1.1	Introduction	1
1.2	Reed-Solomon Codes	3
2	$SL = L$	5
2.1	Zig-zag product	7

1 List decoding of Reed-Solomon codes

1.1 Introduction

An error-correcting code is $C \subseteq \Sigma^n$, where $\Sigma, |\Sigma| = q$ is the alphabet and n is the encoding length.

Define the normalized Hamming distance

$$d(x, y) = \frac{1}{n} |\{i : x_i \neq y_i\}|.$$

We want

$$\delta = \min_{x, y \in C, x \neq y} d(x, y)$$

to be as large as possible (constant as $n \rightarrow \infty$). Imagine balls of radius $\delta/2$ around each point. We can send one of $|C|$ messages by mapping $[|C|] \rightarrow \Sigma^n$; they can withstand $\delta/2$ errors. The rate of the code is $\frac{\lg |C|}{n}$ (how many bits of actual information are sent compared with encoding length); we want the rate and distance to be high, but there are fundamental limits to what can be achieved.

$\delta/2$ is the unique decoding radius. Unique decoding is not possible beyond $\delta/2$.

In many cases, even if there is not a unique codeword within a given radius ($\frac{\delta}{2} + \varepsilon$, say), there may be only a small number of codewords.

Definition 1.1: C is (δ, L) -list decodable if for all $x \in \Sigma^n$, $|B(x, \delta) \cap C| \leq L$.

If you choose your code randomly—that is, choose a codeword, removing a $\frac{\delta}{2}$ ball around it, and repeat—with high probability it will have good list decodability, up to distance $(1 - \varepsilon)\delta$. List decoding is a better way to handle talk about how good the code is than stochastically, because it's worst-case.

The maximum δ is the list-decoding radius. Here $L(n)$ is a function of n . The ideal setting is $L(n) = \text{poly}(n)$. Informally the list decoding radius is the maximum δ such that $L(n) = \text{poly}(n)$.

What does a random code give us; what is possible and not possible? What is the capacity of list decoding?

Theorem 1.2: Let $0 < \delta < 1 - \frac{1}{q}$. (Above this error, the noisy word is essentially random, so we can say nothing.) Then there exists a (δ, L) -list decodable code with rate

$$1 - H_q(\delta) - \frac{1}{L}.$$

Here H_q is defined so that

$$|B(0^n, \delta)| \sim q^{H_q(\delta)n}.$$

Proof. Choose $M = q^{Rn}$ codewords at random. We want

$$\mathbb{P}(C \text{ not } (\delta, L)\text{-list decodable}) < 1.$$

Do a union bound. The probability is at most, by the union bound,

$$\binom{M}{L+1} q^n \left(\frac{|B(x, \delta)|}{q^n} \right)^{L+1} < 1.$$

(q^n is the number of points; the probability of $L+1$ points in a ball centered at a point is $\leq \binom{M}{L+1} \left(\frac{|B(x, \delta)|}{q^n} \right)^{L+1}$.) □

This is almost tight.

Theorem 1.3: If a code is (δ, L) -list decodable with

$$R \geq 1 - H_q(\delta) + \varepsilon$$

then $L(n) \geq q^{\frac{\varepsilon n}{2}}$.

A random ball will contain too many points.

Proof.

$$\begin{aligned} \mathbb{E}_{x \in R\Sigma^n} [|B(x, \delta) \cap C|] &= |C| \frac{|B(0^n, \delta)|}{q^n} \\ &\geq 2^{Rn} q^{n(H_q(\delta)-1)} \\ &\geq q^{\varepsilon n/2}. \end{aligned}$$

□

When $q \rightarrow \infty$, $H_q(\delta) \rightarrow \delta$. We can achieve rate $R \geq 1 - \delta - \varepsilon$ with alphabet of size $2^{O(\frac{1}{\varepsilon})}$ and list size $O(\frac{1}{\varepsilon})$.

Graph: the achievable rates are below the line $\delta + R = 1$.

This is what is known existentially.

The current known explicit codes can achieve $R \geq 1 - \delta - \varepsilon$ with alphabet of size $2^{\text{poly}(\frac{1}{\varepsilon})}$ and list size $n^{O(\frac{1}{\varepsilon})}$. The idea is a folded Reed-Solomon code concatenated with an asymptotically good list-decodable code.

Theorem 1.4 (Johnson bound): Given $C \subseteq \Sigma^n$ with min-distance $\delta = 1 - \varepsilon$, then C is $(1 - \sqrt{\varepsilon}, \text{poly}(n))$ list-decodable.

The Johnson bound says that if you slightly increase the radius, there cannot be an exponential number of codewords.

Theorem 1.5 (Singleton bound): For any code with distance δ , $R \leq 1 - \delta$.

MDS (maximal distance separable) codes are where $R + \delta = 1$, $\delta = 1 - R$, $R = \varepsilon$. MDS codes are $(1 - \sqrt{R}, \text{poly}(N))$ list-decodable.

1.2 Reed-Solomon Codes

We introduce the Reed-Solomon code $RS_{k,n,\mathbb{F}}$, $k < n \leq q$.

These are MDS codes. They achieve $R + \delta = 1$ so are $(1 - \sqrt{R}, \text{poly}(n))$ -list decodable.

Let $\Sigma = \mathbb{F}_q =: \mathbb{F}$. Let

$$C = \{\text{evaluations of degree } k \text{ polynomials on } \{\alpha_1, \dots, \alpha_n\} \subseteq \mathbb{F}\}$$

where $\alpha_1, \dots, \alpha_n$ are distinct and fixed. Codewords correspond to degree k polynomials in $\mathbb{F}[x]$. 2 distinct degree k polynomials can only agree on k points, so $\delta = \frac{n-k}{n} = 1 - \frac{k}{n}$. The rate is $\frac{\log_q(q^{k+1})}{n} = \frac{k+1}{n}$. The Reed-Solomon code is $(1 - \sqrt{\frac{k}{n}}, \text{poly}(n))$ -list decodable.

What is the algorithm? Given a set of points, we need to find all polynomials passing through enough of those points.

Problem 1.6: Given $(\alpha_1, y_1), \dots, (\alpha_n, y_n)$, find all polynomials of degree $\leq k$ such that $p(\alpha_i) = y_i$ for at least \sqrt{nk} indices $i \in [n]$.

Madha Sudan (90's) showed that you can do this with $\sqrt{2nk}$.

Theorem 1.7: There is a polynomial-time algorithm (given in the proof) that given n points as above, finds all degree $\leq k$ polynomials agreeing on $\geq 2\sqrt{nk}$ points.

Proof. The algorithm is as follows. Note this algorithm will work even if the α_i are not distinct.

Define the $(1, k)$ -weighted degree of a polynomial $Q(x, y)$ as $\deg Q(X, Y^k)$. The strategy is as follows.

We will find a low $(1, k)$ -weighted degree polynomial $Q(x, y)$ of degree D such that $Q(\alpha_i, y_i) = 0$ for all i .

Look at $R(x) = Q(x, p(x))$ where $p \in L$. We have

$$\deg R \leq D.$$

For all i such that $p(\alpha_i) = y_i$,

$$R(\alpha_i) = Q(\alpha_i, p(\alpha_i)) = Q(\alpha_i, y_i) = 0.$$

Suppose R has at least t roots and $\deg R \leq D$. If we arrange so that $t > D$, then $R(x) \equiv 0$ identically. Then $Q(x, p(x)) \equiv 0$ implies $y - p(x) \mid Q(x, y)$. (There is a deterministic algorithm to factor bivariate polynomials.) Factor R to find all factors in the form $y - p(x)$; then output those polynomials $p(x)$.

Now we just need to find the polynomial $R(x)$; this is polynomial interpolation.

The number of coefficients in $Q(x, y)$ is $\frac{1}{k} \binom{D+2}{2}$. We need to satisfy n linear constraints; they are linear homogeneous equations in the coefficients. If $\frac{\binom{D+2}{2}}{k} > n$ then there is a nonzero solution: a nonzero $\varphi(x, y)$ of $(1, k)$ weight degree $\leq D$ with $\varphi(\alpha_i) = y_i$ for all i .

If the number of roots is $t > D \approx \sqrt{2kn}$, then we can find all polynomials with agreement t . \square

Guruswami and Sudan improved this to \sqrt{kn} .

Consider $k = 1$, $\mathbb{F} = \mathbb{R}$, find all lines which pass through at least 3 points. If you can interpolate a degree 2 polynomial through all these points, get all lines as lines within the curve. The degree 2 curve will be 2 lines.

Let's look at a small example to see how to improve the bound.

Consider the following picture (see notebook). Here $n = 10$ and $t = 4$. Here we can choose $k = 4$. However it can only have 4 linear factors. The maximum D is 3. Peculiar: through each point there are 2 lines. An algebraic curve passing all points should vanish at the points with multiplicity 2. Fit a polynomial which vanishes with degree 2 at the points.

First define multiplicity.

Definition 1.8: $Q(x, y)$ vanishes with multiplicity r at (α, β) if $Q(x + \alpha, y + \beta)$ doesn't have any monomial of degree $\leq r$.

Lemma 1.9: Let $Q(x, y)$ be with $(1, k)$ degree $\leq D$, vanishing at (α_i, y_i) with multiplicity r for all $i \in [n]$. Let P be a degree k polynomial with agreement $t > D/r$. Then $y - p(x) \mid Q(x, y)$.

Proof. Let $p \in L$, $P(\alpha_i) = y_i$. Define

$$Q^i(x, y) = Q(x + \alpha_i, y + y_i);$$

it has no monomials of degree $< r$. Then

$$\begin{aligned} R(x) &= Q(x, P(x)) \\ &= Q^i(x - \alpha_i, P(x) - y_i) \\ &= Q^i(x - \alpha_i, \underbrace{P(x) - p(\alpha_i)}_{x - \alpha_i \mid P(x) - P(\alpha_i)}). \end{aligned}$$

Q^i has no monomials of degree $< r$. Thus $(x - \alpha_i)^r \mid R(x)$. The number of linear factors is $tr > D$. Thus $R(x) \equiv 0$ and $y - p(x) \mid Q(x, y)$. \square

The number of coefficients in $Q(x, y)$ of $(1, k)$ -degree D is $\frac{1}{k} \binom{D+2}{2}$. The number of homogeneous linear equations is $n \binom{r+1}{2}$. So a nonzero Q exists if $\frac{1}{k} \binom{D+2}{2} > n \binom{r+1}{2}$.

Choose $D = \sqrt{knr(r+1)}$. Let $t = \frac{D}{r} = \sqrt{kn(1 + \frac{1}{r})}$.

Make r large enough, we approach the Johnson bound.

The number of polynomials in the list is at most $\frac{D}{k} = \sqrt{\frac{nr(r+1)}{k}}$, the y -degree.

Conclusion:

1. If $t > \sqrt{kn}$ the list size is $\leq n^\varepsilon$ and we can find all of them.
2. The Reed Solomon code of rate R can be list decoded up to $1 - \sqrt{(1 + \varepsilon)R}$ errors with best size $\frac{1}{\varepsilon\sqrt{R}}$.

This method (the polynomial method) is very flexible. We give another application, the list recovering problem. Given $x \in C$, a noisy channel turns each coordinate into a set s_i and spits out s_1, \dots, s_n , $|s_i| \leq \ell$. We have $|s_i| \leq l$ such that for at least $1 - \delta$ fraction of i 's, $x_i \in s_i$. Find all x such that $x_i \in s_i$ for at least $1 - \delta$ fraction of i 's.

$$\left| \left(\bigcup_{y \in S_1 \times \dots \times S_n} B(y, \delta) \right) \cap C \right| \leq L.$$

Reed-Solomon codes are also good list-recoverable codes.

If $t > \sqrt{kn\ell}$, where t is the number of i such that $x_i \in S_i$, then the Reed-Solomon code is $(1 - \sqrt{k\ell}, \ell, O(n^2\ell^2) = L)$ list recoverable.

We only achieve $1 - \sqrt{R}$. We want to attain $(1 - R, L)$. Guruswami and Rudra came up with folded Reed-Solomon codes. Take a generator γ of \mathbb{F}_q^\times , $n = q - 1$. $\alpha_1, \dots, \alpha_n$ are $S = \{1, \gamma, \dots, \gamma^{q-1}\}$. Consider blocks. \mathbb{F}^m by m . $P(1), \dots, P(r), \dots, P(r^{q-1})$.

It's an open problem to make $n^{O(\frac{1}{\varepsilon})}$ independent of n . Use concatenation to get length down.

You can concatenate a list-recoverable code with a list-decodable codes to get a list-decodable code. $C_{\text{out}} \circ C_{\text{in}}$.

2 $SL = L$

Reingold, Vadhan, and Wigderson.

There are many stories in this problem.

One of the most fundamental questions in theoretical CS is the following.

Claim 2.1: Randomness is useless.

This claim comes from very recent years. 20 years ago people actually believe it's useful (Papadimitriou gave $BPP \neq P$ as a homework exercise in his book).

There are 2 questions: is randomness useful in time and in space?

1. $BPP \stackrel{?}{=} P$
2. $RL \stackrel{?}{=} L$

We focus on the second problem. The first problem is wide open: we don't know $BPP \stackrel{?}{\in} DTIME(2^{o(n)})$. We know more about the second question: Savitch proves

$$BPP \subseteq DSPACE(\ln^2 n).$$

We can do better: Saks and Zhou in 1999 show

$$RL \subseteq DSPACE(\ln^{\frac{3}{2}} n).$$

(They actually show this for BPP.) Think of RL as (undirected) $s-t$ nonconnectivity problem. $coRL$ is interesting because it contains $s-t$ connectivity. (RL doesn't have a complete problem.) For $P \stackrel{?}{=} NP$, just look at a complete problem; here we don't have one.

Here we show $SL = L$. Reingold showed this in 2004.

For simplicity, just think of SL as one problem, undirected $s-t$ connectivity.

Consider a graph.

Jieming starts from a vertex and wants to find a way to Nanjing. Jieming has very little memory, and cannot remember the whole structure of the graph. He can't remember much more than the name of a city.

The standard way to solve this is to take a uniform random walk on the graph. If the graph has $|V| = n$, after $n^2 \ln n$ steps, there is a high probability that he will have visited Nanjing.

But for a general graph, it's necessary to flip $n^2 \ln n$ coins: Consider 2 complete graphs with a bridge: It takes order of n time to hit the bridge vertex, and it has $O(\frac{1}{n})$ chance of crossing the bridge.

For which graphs can we do this randomized routing faster than the worst case? A natural family is the family of expander graphs.

Expander graphs have rapid mixing under a random walk. There are 2 definitions.

Definition 2.2: A graph $G = (V, E)$ is a λ -**edge expander** if for all sets $S \subseteq V$, $|S| \leq \frac{|V|}{2}$,

$$\frac{E(S, \bar{S})}{\text{Vol}(S)} \geq \lambda.$$

Here, $\text{Vol}(S) = \sum_{v \in S} \deg(v)$. (We'll focus on the simple case when the graph is d -regular, i.e., for all $v \in V$ $\deg(v) = d$.) (Pictorially, a subset of $G = (V, E)$ is very spiky, like a sea urchin.)

A graph $G = (V, E)$ is a λ -**spectral expander** if $\lambda_2(L(G)) \geq \lambda$. The Laplacian $L(G)$ is defined from the adjacency matrix A_G . A_G is defined by $(A_G)_{ij} = 1$ if there is an edge between i and j . M is the random walk matrix $\frac{1}{d}A_G$, and

$$L = I - M.$$

(Thinking of this matrix as an operator, $(Lx)_i = \frac{1}{d} \sum_{(i,j) \in E} x_{ij}$, it flows from a vertex to adjacent vertices.)

The Laplacian originates in differential geometry, $L = \text{div } \nabla$.

These 2 definitions are quite close.

Theorem 2.3: Let $h(G)$, $\lambda(G)$ denote the edge and spectral expansions of G . Then

$$\lambda(G) \leq h(G) \leq \sqrt{2\lambda(G)}.$$

The LHS is trivial, the RHS is Cheeger's inequality.

People in CS use a third definition that is more useful.

Definition 2.4: A graph is a λ -expander if for all $x \perp u = \begin{pmatrix} 1 \\ \vdots \\ 1 \end{pmatrix}$, $x \neq 0$,

$$\left| \frac{x^T M x}{x^T x} \right| \leq \lambda.$$

(For edge and spectral expansion we want λ to be large. For random walk expansion we want λ to be close to 0.) $\lambda(L)$ range from 0 to 2, while $\lambda(M)$ range from 1 to -1 .

Note that

$$\left| \frac{x^T (I - L) x}{x^T x} \right| = \left| 1 - \frac{x^T L x}{x^T x} \right|$$

so spectral and random-walk expansion are not quite the same. For bipartite graphs, $\lambda(M) = 1$: random walks are not mixing at all. The side you're on depends on the parity.

For the zig-zag product they consider the third definition.

2.1 Zig-zag product

Think of G_1 having large size and degree (N_1, D_1) , G_2 having small size and degree (N_2, D_2) , and suppose $D_1 = N_2$.

$$G_3 = G_1 \mathbin{\mathbb{Z}} G_2$$

has large size $N_1 N_2$ and small degree, D_2^2 .

Let $\lambda_M(G_i) = \lambda_i$ denote the random walk expansion. Then (a naive bound)

$$\lambda_M(G_3) \leq \lambda_1 + \lambda_2 + \lambda_2^2.$$

It's easy to construct a large degree expander. It's easy to construct small-size expanders (try all possible graphs). It's easy to construct G_1, G_2 , but an expander like G_3 is hard to construct. It takes as input 2 easy-to-construct expanders and outputs a harder to construct graph that's an expander.

Replace each vertex of the original graph with a copy of the second graph G_2 . Its size is hence $N_1 N_2$.

Each vertex in G_3 is labeled $(v, u) \in G_1 \times G_2$. The (i, j) th neighbor is defined as follows. Take u' the i th neighbor of u in G_2 , $u \xrightarrow{i} u'$. Take the u' th neighbor of v in G_1 ; move to the cloud of w ; $v \xrightarrow{u'} w$ in G_2 . Now consider \tilde{u} such that $w \xrightarrow{\tilde{u}} v$ in G_1 ; we move to (w, \tilde{u}) . Now move $\tilde{u} \xrightarrow{j} \tilde{u}'$ in G_2 . Define

$$(v, u) \xrightarrow{(i,j)} (w, \tilde{u}').$$

What's the magic of the zig-zag product? We claim

$$A_3 := A(G_3) = \widetilde{A_2} \widetilde{A_1} \widetilde{A_2}$$

where $\widetilde{B_2} = I \otimes A_2$ has N_1 blocks and each is a copy of A_2 , and $\widetilde{A_1}$ is a permutation (actually matching) matrix where $(\widetilde{A_1})_{((v,u),(w,u'))} = 1$ if

$$v \xrightarrow{u, G_1} w \xrightarrow{u', G_1} v.$$

Consider when $(u, v) \xrightarrow{(i,j)} (w, \tilde{u}')$:

$$\begin{array}{ccc} u & \xrightarrow{i, G_2} & u' \\ v & \xrightarrow{u', G_1} & w \xrightarrow{\tilde{u}, G_1} v \\ \tilde{u} & \xrightarrow{j} & \tilde{u}' \end{array}$$

What is the action of $\widetilde{A_2} \widetilde{A_1} \widetilde{A_2}$ on $e_{(v,u)}$? It goes to $\sum_{uu' \in G_2} e_{(v,u')} = e_v \otimes A_2 e_u$. “Stay in the block of v , move to all the neighbors of u .”

In the second step, we move across blocks as given by $\widetilde{A_1}$.

We need to show

$$\forall x \in \mathbb{R}^{N_1 \times N_2}, x \perp u \implies \frac{x^T M_3 x}{x^T x} \leq \lambda_1 + \lambda_2 + \lambda_2^2.$$

Decompose x as a vector that is uniform on each block,

$$x = \underbrace{\alpha \otimes u}_{\alpha^\parallel} + \underbrace{x'}_{\alpha^\perp}$$

where x' has blocks $\tilde{\alpha}_i \perp u_i$. Then (noting $\widetilde{A_2} \alpha^\parallel = \alpha^\parallel$),

$$\begin{aligned} \langle x, \widetilde{A_2} \widetilde{A_1} \widetilde{A_2} x \rangle &= \langle \widetilde{A_2} x, \widetilde{A_1} \widetilde{A_2} x \rangle \\ &= \langle \widetilde{A_2}(\alpha^\parallel + \alpha^\perp), \widetilde{A_1} \widetilde{A_2}(\alpha^\parallel + \alpha^\perp) \rangle \\ &= \langle \alpha^\parallel + \widetilde{A_2} \alpha^\perp, \widetilde{A_1}(\alpha^\parallel + \widetilde{A_2} \alpha^\perp) \rangle \\ &= \langle \alpha^\parallel, \widetilde{A_1} \alpha^\parallel \rangle + \langle \alpha^\parallel, \widetilde{A_1} \widetilde{A_2} \alpha^\perp \rangle + \langle \widetilde{A_2} \alpha^\perp, \widetilde{A_1} \alpha^\parallel \rangle + \langle \widetilde{A_2} \alpha^\perp, \widetilde{A_1} \widetilde{A_2} \alpha^\perp \rangle \\ &\leq \lambda_1 + 2\lambda_2 + \lambda_2^2 \end{aligned}$$

Because $\alpha_1 + \dots + \alpha_{N_1} = 0$ we can use the expansion properties of G_1 to get the λ_1 bound for the first term.

The idea:

1. Suppose G has parameters (N, D) and expansion λ . (say $1 - \frac{1}{ND}$)
2. G^2 has parameters (N, D^2) , expansion λ^2 .

3. Take H with parameters $(D^2, \sqrt{D}), \lambda_2$. Then

$$G^2 \mathbb{Z} H = G_{\text{new}}$$

has parameters (ND^2, D) expansion $\lambda(G_{\text{new}}) \leq \lambda^2 + 2\lambda_2 + \lambda_2^2 \leq \lambda^2 + \varepsilon$.

We go from G with parameters (N, D, λ) to G_{new} with parameters $(ND, D, \lambda^2 + \varepsilon)$. Call this operator Z . Doing this operator t times, $Z^t G = (ND^{2^t}, D, \lambda^{2^t} + \varepsilon')$. Then set $t = \lg\left(\frac{ND}{2}\right)$ to get $Z^t G$ with parameters $(N^2 D, D, \frac{1}{e} + \varepsilon)$. (We're cheating a little, using the naive bound. We need a better bound to get ε small: G_3 has expansion $\frac{1}{2}(1 - \lambda_2^2)\lambda_1 + \frac{1}{2}\sqrt{(1 - \lambda_2^2)^2\lambda_1^2 + t\lambda_2^2}$. To go from our proof to the real proof, you just need to note $\alpha^{\parallel} \perp \alpha^{\perp}$, $\alpha^{\parallel} \perp \widetilde{A_2}\alpha^{\perp}$. Plug in this picture into the formula.