

# 1 Backpropagation

<http://arxiv.org/pdf/1411.6191.pdf>

We'll go through backpropagation in too much detail.

1. It's 30 years old.
2. It's the workhorse of deep learning.
3. It's a large contributor to state-of-the-art.
4. It's dead simple: gradient descent (computing derivatives) plus the chain rule.
5. But the brain doesn't use it. Neuroscientists have tried for 50 years to find evidence of backprop in the brain and are pretty convinced it's not there.

Backprop is more complicated than the brain; it has feedforward and feedbackward output. The brain seems to be working better. So we take another look and ask, can we simplify backprop?

## 1.1 Learning algorithm with linear outputs

Given a weight vector  $\mathbf{w}$ , an input  $\mathbf{x}$ , output  $\langle \mathbf{w}, \mathbf{x} \rangle$ . We are given a loss function, for example the **mean squared error**

$$\ell(\langle \mathbf{w}, \mathbf{x} \rangle, y) = \frac{1}{2}(y - \langle \mathbf{w}, \mathbf{x} \rangle)^2.$$

Minimizing it is doing linear regression. Other loss functions are logistic and linear loss. This is studied most in online learning (Elad Hazan). We'll use MSE because it's simplest.

We won't use a neural network to minimize linear loss, but the quantity naturally comes up.

We need a learning rule.

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \eta \nabla_{\mathbf{w}} \ell_{\mathbf{w}}.$$

If loss is convex, this converges.

The only problem is that the class of functions is very restricted, so we would like to extend this to more complicated outputs than just dot products. The standard answer is to build a neural network.

Consider a neural network. The hidden nodes (rectilinear/rectifiers) have the following form

$$S_{\mathbf{w}^j}(\mathbf{x}) = \max\{0, \langle \mathbf{w}^j, \mathbf{x} \rangle\}.$$

This is the closest we can get to linear. If we didn't have nonlinear functions, we'd just get linear functions, and there would be no advantage to a network. Leshno showed that  $\text{span}\{S_w(x) : w \in \mathbb{R}^d\}$  is dense in  $\mathbb{R}^d$ . Assume the output layer is linear (but it doesn't matter). We have a loss  $\{\text{output loss}\} \times \{\text{labels}\} \rightarrow \mathbb{R}$ .

We have a function  $F_{\mathbf{w}}(x)$  computed by the neural net. To update  $\mathbf{w}$ , set

$$\mathbf{w}_{t+1}^j = \mathbf{w}_t^j + \eta \nabla_{\mathbf{w}^j} \ell(F_{\mathbf{w}^j}(x), \text{label}).$$

Here  $F_{\mathbf{w}}$  is not convex so we're not guaranteed to find a global minimum, but it's good enough.

We'll deconstruct backprop and find hidden degrees of freedom.

Applying the chain rule,

$$\nabla_{\mathbf{w}^j} \ell(F_{\mathbf{w}^j}(x), \text{label}) = \underbrace{\nabla_{\mathbf{x}^{\text{out}}} \ell(\mathbf{x}^{\text{out}}, \text{label})}_{1 \times d} \cdot \underbrace{(\nabla_{\mathbf{x}^j} x^{\text{out}})}_{d \times 1} \cdot \underbrace{(\nabla_{\mathbf{w}^j} x^j)}_{n_j}$$

The scalar that's the product of the first 2 vectors is written  $\delta^j$ . Our goal is to simplify this.

We compute  $\nabla_{x^j} \mathbf{x}^{\text{out}} = ?$ . Compute derivative of  $S_{\mathbf{w}^j}(\mathbf{x}) = \max(0, \langle \mathbf{w}, \mathbf{x} \rangle)$ ,

$$\frac{\partial S_{\mathbf{w}}(x)}{\partial x_L} = \begin{cases} w_i, & \langle \mathbf{w}, \mathbf{x} \rangle > 0 \\ 0, & \text{else.} \end{cases} = w_i \mathbb{1}_j$$

where  $\mathbb{1}_j$  is 1 if node  $j$  is active and 0 else.

Hence  $\nabla_{\mathbf{x}} S_{\mathbf{w}^j}(\mathbf{x}) = \mathbf{w} \mathbb{1}_j$ . We get

$$\nabla_{x^j} \mathbf{x}^{\text{out}} = \sum_{\{k: j \rightarrow k\}} w^{jk} \mathbb{1}_k \left( \sum_{\{l: k \rightarrow l\}} w^{kl} \mathbb{1}_l \left( \sum \dots \right) \right) = \sum_{\substack{\text{active paths} \\ \text{from } j \text{ to } x^{\text{out}}}} (\text{weights on path}).$$

Go through every layer, add the weights if they acted.

Define  $\pi^j = \nabla_{x^j} \mathbf{x}^{\text{out}} = \sum_{\text{active paths}} (\text{weights})$ , the influence.

We don't care about  $\nabla \ell$ ; this is the business end (interaction of neural network with data). We care about the structure of the neural network,  $\nabla_{x^j} \mathbf{x}^{\text{out}}$ , how the neural network communicates with itself. This is difficult to do in hardware. Why? You have a system with 2 regimes: input flows forwards, then backpropagate error. Feedforward phase and feedback phase, with a global clock. You need massive synchrony. It's not realistic to build an organism which does this.

It's hard to simplify derivatives and chain rule. We need a new angle. Instead of thinking about it as a giant organism, we'll see each node has its own loss and is doing a learning algorithm.

## 1.2 Each node is a learning algorithm

We have

$$\nabla_{\mathbf{w}^j} x^j = \mathbf{x} \mathbb{1}_j$$

(recall  $x^j = S_{\mathbf{w}^j}(\mathbf{x}) = \max(0, \langle \mathbf{w}^j, \mathbf{x} \rangle)$ ).

Let's take a step back.

Consider backpropagation for 1 node. It does the following.

1. Receive input.

2. Output.
3. Receive backpropagated error.
4. Update weights.

Step	Quantity
Receive input	$\phi^j$
Output	$\max\{0, \langle \mathbf{w}^j, \phi^j \rangle\}$
backpropagated error	$\delta^j$
update weights	$\mathbf{w}_{t+1}^j = \mathbf{w}_t^j + \eta^j \delta^j \phi^j \mathbb{1}^j$

Define the rectilinear loss

$$\ell_{\text{RL}}(\mathbf{w}, \phi, \delta) = \begin{cases} \delta \langle \mathbf{w}, \phi \rangle, & \text{if node is active} \\ 0, & \text{else.} \end{cases}$$

This is not convex but it is selectively convex. Every node either does nothing  $\langle \mathbf{w}^j, \phi^j \rangle < 0$  or minimizes a linear loss using gradient descent.

This lets you prove error and generalization loss, even in adversarial settings. We get nice generalization bounds on the nodes.

How can we use this to simplify backprop?  $\delta$  is computed by backprop. We can imagine replacing  $\delta$  by another number. Let's mess with  $\delta$ ; replace it with something easier to compute. (Obviously you have to be careful, as  $\delta$  was chosen to minimize global loss.)

Consider the special case where the output layer has 1 node and

$$\nabla_{w^j} \ell(NN, \text{label}) = \nabla_{x^{\text{out}}} \left( \sum (\text{weights on active paths}) \right) (\nabla_{\mathbf{w}^j} x^j).$$

The first 2 quantities are scalars; multiplying gives  $\delta^j$ .

We'll do something naive which works.

$\pi^j = \sum_{\text{all active paths}}$ ,  $\tau^j = \sum_{\text{all active paths of length 1}} (\text{weight on path}) = \sum_{\{k:j \rightarrow k\}} w^{jk} \mathbb{1}_k$ . As the network gets bigger,  $\tau^j$  gets easier and easier relative to  $\pi^j$ .

We define a new algorithm called kickback.

$$\begin{aligned} \text{BP} &= \ell_{\text{RL}}(\mathbf{w}^j, \phi^j, \delta) \\ \text{KB} &= \ell_{\text{RL}}(\mathbf{w}^j, \phi^j, (\nabla_{\mathbf{x}^{\text{out}}} (\text{Poss})) \tau^j). \end{aligned}$$

The analogue in biology is neural modulators. They are diffuse signals received by a large part of the brain.

We have a global scalar,  $\nabla_{x^{\text{out}}} (\text{loss})$ . Node  $j$  just looks at adjacent nodes, and asks which one was active after I was active; sum up the weights to compute  $\tau$ .  $\tau^j$  is easy to compute because it's completely local information. Plus we get global information: whether we overshot or undershot. This is biologically plausible.

Kickback is no longer gradient descent, so we have no guarantee of error reduction. There is an easy way to fix this.

**Definition 1.1:** A node is **coherent** if  $\tau^j > 0$ .

A network is **coherent** if  $\tau_j > 0$  for all  $j$ .

We need a condition for consistency: the effect of a neuron is always positive or always negative. All the positive rectifiers have positive weights and all the negative rectifiers have negative weights. This is a property of the network. The price you pay for throwing away backprop is constraint on the network.

Take a NN, initialize weights randomly, change the weights to be consistent, and forgot about it, without worrying about coherency. It worked.

Conclusion: Coherent initialization and kickback give results almost identical to backprop. RPROP was invented in 90's to speed up optimization; RMS-PROP is a more complicated version with many batches. Compute  $\delta^j$ , and throw away everything except for the sign. It's a hack that works surprisingly well. About half of deep learning papers rely on RMS-PROP.

Coherency: Starting it at 1, it typically went down to .9; backprop didn't seem to care (it moved everywhere). Backprop both doesn't need and doesn't preserve coherence. Kickback needs and seems to preserve coherence. When it's not initialized for coherence, often it doesn't work (blows up), or does weird oscillations.

Everything about kickback is simpler than the brain except coherency. Excitatory neurons behave as they should, but inhibitory neurons don't work the way they should; if we can fix this then we have a plausible biological algorithm.

This only works with binary classification. Multiple classification is hard. "I'm interested in reinforcement learning, excited about TD-error, a scalar."

Mixture version of gradient descent. Each part on one computer? Kickback is more parallelizable.

It's not faster (only marginally faster) in software; it's faster in hardware.

## 2

Preparation

- Use softmax regression for multi-class classification.
- Weight decay (regularization).

Autoencoders do dimension reduction.

### 2.1 Intriguing properties of neural networks

- Activation of a hidden unit is a meaningful feature. Try to find  $x' = \arg_{x \in \text{image set}} \max \langle \phi(x), e_i \rangle$ . But randomly choosing a vector  $v$

Adversarial example. Why do they exist? The mappings are discontinuous to a significant extend. Cross model generalization: Adversarial examples trained against A may make (different architecture) B fail. Cross training-set generalization: even if B trained with different data.

Minimize  $\|r\|_2$  subject to  $f(x + r) = l$ .

Adversarial examples are universal. Back-feeding may improve generalization of model.

Are the adversarial examples robust? Are the adversarial examples islands, or do they have a lot of area around them? How much Gaussian noise to get same error as adversarial perturbation?

The opposite direction: unrecognizable images that make the DNN believe with high confidence.

Use evolutionary algorithms. Multidimensional: if the prediction score is higher than current highest score of any class, keep it. (Each class is a niche.)

Indirect encoding: regular images with meaningful patterns. Compositional pattern-producing network. Images artificial for humans but NN can recognize it.

Have another category, “none of the above.”

Regular images that fool MNIST network. ex. 1 has vertical strokes.

ImageNet: most classes aren’t fooled, a few are. Indirect encoding also give high confidence. (It fails for dogs and cats because there are too many classes for cats and dogs.)

ImageNet: regular images.

Fool closely related classes. Different runs of the algorithm give completely different images: there are many ways to fool the DNN.

Repetition of patterns increases confidence. DNN learn low/mid-level patterns rather than global structure. Many natural images do have multiple copies.

3. Train with fooled images help it resist? Retraining does not help in terms of worst-case adversarially. **But is the probability smaller?**

The problem is the way we use the neural nets?

4. Why do adversarial examples exist? Extreme nonlinearity, insufficient model averaging/regularization.

But linear behavior in high-dimensional spaces is sufficient to cause adversarial examples.

Perturbation is meaningless if  $\|\eta\|_\infty < \varepsilon$ . Adversarial example  $\tilde{x} = x + \eta$ . Get  $w^T \eta$ . Increase of activation is  $\varepsilon mn$ , activation increases as dimension. (I.e.,  $\|\cdot\|_2^2 \leq n \|\cdot\|_\infty$ .)

Adversarial training of linear models, deep nets helps.

Why do they generalize? When different models misclassify, they often agree with each other.

Hypothesis: neural networks trained all resemble the linear classifier learned on the same training set.

Just one type of adversarial example?

Note: ImageNet much better than MNIST.

Just use gaussian fooling. Previous papers overkill.

### 3 Neural Turing machines

Motivation: most machine learning algorithms try to learn a static mapping and it has been elusive to incorporate memory in the learning.

Before Turing invented Turing machines, they build circuits without memory. Turing realized that memory can encode the state; it’s more powerful than a functional mapping. You can do more complicated stuff. A Turing machine works for all sizes of inputs.

3 papers:

1. Learning to execute (recursive neural net). (NN encodes in hidden states.)
2. QA memory network.
3. Neural Turing machine.

Unlike CNN, in RNN the output relies on the hidden state of the previous time stem. LSTM is a special case of RNN that is made to store long term memory easily.

1. Learning to execute: Can LSTM learn to execute python code?

Operators: addition, subtraction, multiplication, variable assignments, if statements, and for loops but not double loops.

Length parameter: constrain the integer in a maximal length. Nesting parameter: constrain the number of times to combine operations. **How did they generate the programs? ‘(’ is 1 character; “if” is one unit? Only binary operations.**

They use **curriculum learning** (cf. Winston’s arch learning): gradually increase the difficulties of training examples. 4 methods:

- (a) baseline: training examples with length  $a$  and nesting  $b$ .
- (b) naive: gradually increase.
- (c) mix: pick random lengths and nesting.
- (d) combined: combination of naive and mix.

Teacher forcing: when predicting the  $i$ th digit, provide it with correct  $i - 1$  digits. **I think the experimenters did too much “feeding.”**

2. QA memory networks: Instead of using a recurrent matrix to retain information through time, build a memory directly. Train the model to operate effectively with the memory component.

Framework:

- (a) input feature map
- (b) generalization
- (c) output feature map
- (d) response.

Learn the parameters of  $U_O, U_R$  where the scoring function is  $S(x, y) = \Phi(x)^T U^T U \Phi(y)$ , using gradient descent.

Knowledge database only, interaction between computation resources and memory is limited.

### 3. NTM. Attention as parameters.

$$\begin{aligned} \text{Read: } \|w_t\|_1 &= 1 \\ r_t &\leftarrow w_t \cdot M_t \\ \text{Write: } \widetilde{M}_t &\leftarrow M_{t-1} \cdot (1 - e_t w_t) \\ M_t &\leftarrow \widetilde{M}_t + a_t w_t. \end{aligned}$$

Addressing mechanisms

- (a) content-based: cosine similarity. (weight by  $e^{\beta_t K(k_t, M_t(i))}$ )
- (b) interpolation
- (c) shifting and sharpening **shift attention?** Weightings get soft so sharpen them.

Only fuzziness from convolution? 3 modes

- (a) weights chosen by content system without modification of location system
- (b) weights from content system can be chosen and shifted
- (c) weights from the previous time step can be rotated without any input from content addressed. Allows iteration.

How to model controller network: FNN, RNN (better: mix info across time. controller is the CPU, memory is RAM, hidden states are CPU).

Output is not binary.

**Stopping time is hardwired; otherwise there would cause a discontinuity. (Alternative - have fixed zero probability of leaving stopping time?)**

Tasks

- (a) copy
- (b) repeated copy (fails to figure out where to end. use another memory location to help switch back the pointer to the start). **should teach it to count first! can we do subroutines?**
- (c) associative recall (given sequence AB...A, return A). (method: write something to memory afterwards. network do content-based lookup, shift memory by 1.)
- (d) priority sort (write weightings, read weightings) **content-based important for last 2?**

**coerce to Turing machine without weights? go backwards? on this computational mode...**

parse video?